

SMECY C (or SME-C? ☺)

C99 with pragma and API for parallel execution,
processor mapping and communication generation

`Remi.Barrere@thalesgroup.com`

`Ronan.Keryell@hpc-project.com`

May 28, 2011

1 Introduction

In the SMECY project we want to use C source code as a portable intermediate representation (IR) between tools from high-level tools down to lower-level tools because of its good trade-off between expressiveness and readability, without compromising portability.

The targets envisioned in SMECY are heterogeneous multicore systems with shared memory or not, with various hardware accelerator, such as ASIC, ASIP, FPGA, GPU, specialized vector processors, partially reconfigurable accelerators...

Unfortunately, since it is undecidable to get high-level properties from such programs, we use decorations to help tools to understand program behaviour and generate codes for some hardware targets. We try to keep clear decorations, easy to understand, so that SMECY C can also be used as a programming language.

A SMECY program contains various functions that may be executed on various processors, accelerators, GPU... that may consume and produce data from different physical memory spaces.

Since we want to express also performance on given platforms, we keep the opportunity to have platform-specific pragma and API or specialized intrinsic types and functions, for example to express use of special hardware accelerator functions or operations.

The hardware specific pragma and intrinsics are to be defined between software and hardware partners involved in various use cases. But it may not be possible to address all the programming models and platforms envisioned in the project.

1.1 SMECY-C programming model

The programming model is based on C processes, with a virtual shared memory and threads *à la* OpenMP. Since we may have quite asynchronous processes in a real application description or at the back-end level in the execution model, we can cope indeed with different C processes communicating with an API.

We add mapping information stating on which hardware part a function is to be placed and run.

The programming model exposed in the following is based on an OpenMP SMP model because of its (rather) simple readability, elegance, old background and wide acceptance. We make the hypothesis that a SMECY program is a correct OpenMP program that can be executed in sequential with a C OpenMP-free compiler (just by ignoring OpenMP `#pragma`) and in parallel on a SMP target (such an *x86* machine) by using an OpenMP compiler *with the same semantics*. Since we use C (by opposition to a DSL) as an internal representation, we choose this behaviour to stick to standard behaviour as much as possible. This is known as the sequential equivalence. Since we can cope different results for performance reasons from executing in parallel non associative floating point operations, we deal with only *weak* sequential equivalence instead of *strong* sequential equivalence.

Of course, this model is incompatible with a real hardware target envisioned in SMECY, so we need to add hints in the code explaining memory dependencies at the function call levels. Since it is quite difficult to describe general dependencies, we approximate memory dependencies with rectangles (and more generally hyperparallelepiped in any dimension) that can be read, written or both. We think these abstractions are good trade-offs between expressiveness (and what a programmer can endure...) and hardware capabilities. Even if there is some similitude with HPF (High Performance Fortran) we do not deal with strides¹.

With this information the tools can guess the communication to generate between the different functions and memory spaces to emulate the global OpenMP memory semantics.

The neat side effect is that we have the same global program executed on all the platforms (sequential, real OpenMP and SMECY) with the same semantics and we can see the sequential version as a functional simulator of the SMECY application and the real OpenMP version as a parallelized quicker version of this simulator.

It is also easy to debug the application, but also all the SMECY tools used or developed in the project.

To be able to address real hardware from the C level with special needs:

- to specify hardware register names;
- define input/output routines specifically but in portable way;
- define fixed point computations;
- specific data size;
- accumulator register (DSP...);
- different memory spaces that can be chosen specifically to optimize storage and speed (DSP, hardware accelerators with scratch pad memory...);
- saturated arithmetic.

¹Indeed, by using some higher dimensional arrays than the arrays used in the application, you can express them... So may be we can express them in the syntax?

For all these we rely on the TR 18037 Embedded C standard, supported for example by ACE tools.

To describe different processes communicating together in an asynchronous way, we do not have anymore a sequential equivalence and then do not use pragma to express this. So we use a simple communication, synchronization and threading API. Since we target embedded systems with a light efficient implementation, we can rely on such standard API as the ones of the MultiCore Association: MCAPI (communication), MRAPI (synchronization) and MTAPI (threading).

As modern programs need a clear documentation, such as with Doxygen marking style providing meta information on different elements of the program, we can use this (hopefully correct) information to help the compiling process itself.

Since tools are to be oriented for a specific target, taking into account various hardware and compilation parameters that are to be kept orthogonally to application sources, those parameters are kept aside in some description files that flow between tools. These files may be represented with an XML syntax (using a SMECY naming space) or by even simpler format (JSON, YAML...).

1.2 Reference documents

Besides the SMECY documentation, the reader should be knowledgeable of some work of the ISO/IEC JTC1/SC22/WG14 committee on the C language standard:

- ISO/IEC 9899 - Programming languages - C (Technical Corrigendum 3 Cor. 3:2007(E)) <http://www.open-std.org/jtc1/sc22/wg14/www/docs/n1256.pdf>
- TR 18037: Embedded C <http://www.open-std.org/jtc1/sc22/wg14/www/docs/n1169.pdf>
- Future C1X standard <http://www.open-std.org/jtc1/sc22/wg14/www/docs/n1548.pdf>

and other standards such as

- MCA (MultiCore Association) API: MCAPI, MRAPI & MTAPI, www.multicore-association.org
- HPF (High Performance Fortran) <http://hpff.rice.edu/versions/hpf2>
- ISO/IEC TR 24717:2009, Information technology – Programming languages, their environments and system software interfaces – Collection classes for programming language COBOL 20xx <http://www.cobolstandard.info/j4/files/std.zip>

2 Intermediate representation use cases

2.1 Direct programming

A programmer can program her application directly in SMECY C with OpenMP and SMECY-specific pragma and API to target a SMECY platform.

It can be at rather high-level, by using only high-level pragma, or rather at a lower level, by using different communicating processes with the API and even specialized API and pragma for specific hardware.

A process written in C code with pragma and API can express a global host controlling process of an application or a local program in a specialized processor. And we may have many of such processes to express different producer and consumer Kahn processes interacting through a NoC in an asynchronous way.

2.2 System high-level synthesis

A compiler can take a sequential plain C or Matlab or another language code, analyze and parallelize the code by adding automatically parallel and mapping pragma. This can be seen as a high-level synthesis at system level.

A tool such as Par4All can do these kinds of transformation.

2.3 Hardware high-level synthesis

A compiler can take a program with SMECY pragma and compile any call with a mapping of a given kind into some hardware configuration or program to be executed instead of the function and an API call to use this hardware part from the host program.

Since the pragma are designed to be concretely compilable, such a tool should be easy to do with a simple compilation framework, such as ROSE Compiler.

The SMECY API and intrinsics are chosen to be mapped quite straightforwardly to real hardware functions by the back-end.

3 Exemples

3.1 Program with contiguous memory transfers

During C memory transfer, if we work on arrays with the last dimension taken as a whole, the memory is contiguous and the programs often work even there are some aliasing such as using a 2D array zone as a linearized 1D vector.

The following program exposes this kind of code where some work sharing is done by contiguous memory blocks.

```

1  /* To compile this program on Linux, try:

      make CFLAGS='-std=c99 -Wall' pragma-example

   To run:
6  ./pragma-example; echo $?
   It should print 0 if OK.

   You can even compile it to run on multicore SMP for free with

11 make CFLAGS='-std=c99 -fopenmp -Wall' pragma-example

   To verify there are really some clone() system calls that create the threads:
   strace -f ./pragma-example ; echo $?
```

```

16      You can notice that the #pragma smecy are ignored (the project is
      on-going :-)) but that the program produces already correct results in
      sequential execution and parallel OpenMP execution.

      Enjoy!

21      Remi.Barrere@thalesgroup.com
      Ronan.Keryell@hpc-project.com
      for ARTEMIS SMECY European project.
    */
26  #include <stdbool.h>

    /* function Gen

31      Example of old C89 array use-case where the size is unknown. Note that
      this implies some nasty access linearization with array with more than
      1 dimension.
    */
    void Gen(int *out, int size) {
36      // Can be executed in parallel
    #pragma omp parallel for
      for (int i = 0; i < size; i++)
        out[i] = 0;
41    }

    /* function Add

      Nice C99 array with dynamic size definition. Note this implies having
46      array size given first
    */
    void Add(int size, int in[size], int out[size]) {
      // Can be executed in parallel
    #pragma omp parallel for
51      for (int i = 0; i < size; i++)
        out[i] = in[i] + 1;
    }

56  /* function Test */
    bool Test(int size, int in[size]) {
      bool ok = true;
      /* Can be executed in parallel, ok is initialized from global value and
      at loop exit ok is the  $\&\&$  operation between all the local ok
61      instances: */
    #pragma omp parallel for reduction(&&:ok)
      for (int i = 0; i < size; i++)
        /* We cannot have this simple code here:
          if (in[i] != 2)
66          exit(-1);
          because a loop or a fonction with exit() cannot be executed in parallel.

      Proof: there is a parallel execution interleaving that may execute

```

```

71      some computations in some threads with a greater i that the one
       executing the exit() done on another thread. So the causality is
       not respected.

       Anyway, in an heterogenous execution, just think about how to
       implement the exit() operating system call from an
76      accelerator... No hope. :-)

       So use a reduction instead and return the status for later
       inspection:

      */
81      ok &= (in[i] == 2);

      // Return false if at least one in[i] is not equal to 2:
      return ok;
    }
86

    /* main */
    int main(int argc, char* argv[]) {
      int tab[6][200];
91      // Gen is mapped on GPP 0, it produced (out) an array written to arg 1:
      #pragma smecy map(GPP, 0) arg(1, [6][200], out)
      /* Note there is an array linearization here, since we give a 2D array
       to Gen() that uses it. This is bad programming style, but it is just
       to show it can be handled in the model :-) */
96      Gen((int *) tab, 200*6);

      // Launch different things in parallel:
      #pragma omp parallel sections
      {
101      // Do one thing in parallel...
      #pragma omp section
      {
        /* Map this "Add" call to PE 0, arg 2 is communicated as input as an
         array of "int [3][200]", and after execution arg 3 is
106        communicated out as an array of "int [3][200]"

         Note the aliasing of the 2 last arguments. Just to show we can
         handle it. :*/
        #pragma smecy map(PE, 0) arg(2, [3][200], in) arg(3, [3][200], out)
111        Add(200*3, (int *) tab, (int *) tab);
      }
      // ...with another thing
      #pragma omp section
      {
116        /* Map this "Add" call to PE 1, arg 2 is communicated as input as an
         array of "int [3][200]" from address tab[3][0], that is the
         second half of tab, and after execution arg 3 is communicated out
         as an array of "int [3][200]", that is the second half of tab

         Note the aliasing of the 2 last arguments. Just to show we can
121        handle it. :*/
        #pragma smecy map(PE, 1) arg(2, [3][200], in) \

```

```

                                arg(3, [3][200], out)
    Add(200*3, &tab[3][0], &tab[3][0]);
126     }
    }

    // Launch different things in parallel:
    #pragma omp parallel sections
131     {
        #pragma omp section
        {
            #pragma smecy map(PE, 2) arg(2, [2][200], in) arg(3, [2][200], out)
            Add(200*2, (int *) tab, (int *) tab);
136        }
        #pragma omp section
        {
            #pragma smecy map(PE, 3) arg(2, [2][200], in) arg(3, [2][200], out)
            Add(200*2, &tab[2][0], &tab[2][0]);
141        }
        #pragma omp section
        {
            #pragma smecy map(PE, 4) arg(2, [2][200], in) arg(3, [2][200], out)
            Add(200*2, &tab[4][0], &tab[4][0]);
146        }
    }
    // An example where arg 2 is just used as a whole implicitly:
    #pragma smecy map(GPP, 0) arg(2, in)
    bool result = Test(200*6, (int *) tab);
151    // Return non 0 if the computation went wrong:
    return !result;
}

```

3.2 Program with non-contiguous memory transfers

In the following example, we apply different computations on square pieces of the image, that do not have contiguous representation in memory. That is why we need to express restrictions on the use of the whole array.

```

/* To compile this program on Linux, try:
2
    make CFLAGS='-std=c99 -Wall' example_2D

    To run:
    ./example_2D; echo $?
7    It should print 0 if OK.

    You can even compile it to run on multicore SMP for free with

    make CFLAGS='-std=c99 -fopenmp -Wall' example_2D
12

    To verify there are really some clone() system calls that create the threads:
    strace -f ./example_2D ; echo $?

    You can notice that the #pragma smecy are ignored (the project is
17    on-going :-)) but that the program produces already correct results in

```

```

        sequential execution and parallel OpenMP execution.

        Enjoy!

22     Ronan.Keryell@hpc-project.com
        for ARTEMIS SMECY European project.
    */

#include <stdlib.h>
27 #include "example_helper.h"

// Problem size
enum { WIDTH = 500, HEIGHT = 200 };
32

/* The main host program controlling and representing the whole
   application */
int main(int argc, char* argv[]) {
37     int image[HEIGHT][WIDTH];
    unsigned char output[HEIGHT][WIDTH];

    // Initialize with some values
    init_image(WIDTH, HEIGHT, image);
42

    #pragma omp parallel sections
    {
        // On one processor
        // We rewrite a small part of image:
47     #pragma smecy map(PE, 0) arg(3, inout, [HEIGHT][WIDTH]
                                   /[HEIGHT/3:HEIGHT/3 + HEIGHT/2 - 1] \
                                   [WIDTH/8:WIDTH/8 + HEIGHT/2 - 1])
        square_symmetry(WIDTH, HEIGHT, image, HEIGHT/2, WIDTH/8, HEIGHT/3);

52     // On another processor
    #pragma omp section
        // Here let the compiler to guess the array size
    #pragma smecy map(PE, 1) arg(3, inout, /[HEIGHT/4:HEIGHT/4 + HEIGHT/2 - 1] \
                                   [3*WIDTH/8:3*WIDTH/8 + HEIGHT/2 - 1])
57     square_symmetry(WIDTH, HEIGHT, image, HEIGHT/2, 3*WIDTH/4, HEIGHT/4);

        // On another processor
    #pragma omp section
        // Here let the compiler to guess the array size
62     #pragma smecy map(PE, 1) arg(3, inout, /[2*HEIGHT/5:2*HEIGHT/5 + HEIGHT/2 - 1] \
                                   [WIDTH/2:WIDTH/2 + HEIGHT/2 - 1])
        square_symmetry(WIDTH, HEIGHT, image, HEIGHT/2, WIDTH/2, 2*HEIGHT/5);
    }
    // Here there is a synchronization because of the parallel part end
67

    // Since there
    normalize_to_char(WIDTH, HEIGHT, image, output);

    write_pgm_image("2D.example-output.pgm", WIDTH, HEIGHT, output);

```



```

72     return EXIT_SUCCESS;
    }

```

3.3 Pipelined example

TODO

3.4 Remapping example

Some information can be in a given layout but needed in another layout to be used by a specific hardware accelerator.

```

1  #include <stdlib.h>
    #include "example_helper.h"

    // Problem size
    enum { WIDTH = 500, HEIGHT = 200, LINE_SIZE = 100 };

6  /* Apply some pixel value inversion in a 1D array
    */
    void
    invert_vector(int line_size,
11         int input_line[line_size],
            int output_line[line_size]) {
        for(int i = 0; i < line_size; i++)
            output_line[i] = 500 - input_line[i];
    }

16

    /* The main host program controlling and representing the whole
       application */
    int main(int argc, char* argv[]) {
21     int image[HEIGHT][WIDTH];
        unsigned char output[HEIGHT][WIDTH];

        // Initialize with some values
        init_image(WIDTH, HEIGHT, image);

26     // Draw 70 horizontal lines and map operation on 8 PEs:
    #pragma omp parallel for num_threads(8)
        for(int proc = 0; proc < 70; proc++)
            // Each iteration is on a different PE in parallel:
31     #pragma smecy map(PE, proc & 7) \
                arg(2, in, [1][LINE_SIZE]) \
                arg(3, out, [1][LINE_SIZE])
            // Invert an horizontal line:
            invert_vector(LINE_SIZE,
36                &image[HEIGHT - 20 - proc][WIDTH/2 + 2*proc],
                &image[HEIGHT - 20 - proc][WIDTH/2 + 2*proc]);

        /* Here we guess we have 5 hardware accelerators and we launch
           operations on them: */
41     #pragma omp parallel for num_threads(5)

```

```

    for(int proc = 0; proc < 5; proc++) {
        /* This is need to express the fact that our accelerator only accept
           continuous data but we want apply them on non contiguous data in
           the array */
46     int input_line[LINE_SIZE];
        int output_line[LINE_SIZE];
        /* We need to remap data in the good shape. The compiler should use
           the remapping information to generate DMA transfer for example and
           remove input_line array */
51     SMECY_remap_int2D_to_int1D(HEIGHT, WIDTH, HEIGHT/3, 30 + 20*proc,
                                LINE_SIZE, 1, image,
                                LINE_SIZE, input_line);
        // Each iteration is on a different PE in parallel:
    #pragma smecy map(PE, proc) arg(2, in, [LINE_SIZE]) arg(3, out, [LINE_SIZE])
56     invert_vector(LINE_SIZE, input_line, output_line);
        SMECY_remap_int1D_to_int2D(LINE_SIZE, output_line,
                                    HEIGHT, WIDTH, HEIGHT/3, 30 + 20*proc,
                                    LINE_SIZE, 1, image);
    }
61
    // Convert int image to char image:
    normalize_to_char(WIDTH, HEIGHT, image, output);

    write_pgm_image("remapping-example-output.pgm", WIDTH, HEIGHT, output);
66
    return EXIT_SUCCESS;
}

```

4 SMECY embedded C language

We take as input C99 (ISO/IEC 9899:2007) language with extension for embedded systems (TR 18037).

Refer to these document for more information.

5 Description of the SMECY directives

The generic format of SMECY code decorations are language dependent, because if here we describe a SMECY IR implementation based on C, it is indeed more general.

- In C/C++:

```
#pragma smecy clause[[,]clause]... newline
```

We can use \ at the end of line for continuation information.

- In Fortran:

```
!$smecy clause[[,]clause]... newline
```

Use & at the end of line for continuation information.

- In other languages: use **#pragma** equivalent, if not available, use comments *à la* Fortran. For example in Python:

```
#$smecy clause[[,clause]... newline
```

Use also **&** at the end of line for continuation information

In implementations that support a preprocessor, the **_SMECY** macro name is defined to have the decimal value *yyyymm* where *yyyy* and *mm* are the year and month designations of the version of the SMECY API that the implementation supports. If this macro is the subject of a **#define** or a **#undef** preprocessing directive, the behavior is unspecified.

5.1 OpenMP support

SMECY is based on OpenMP. It is not clear yet what level of OpenMP is supported. Since a SMECY platform is made at least from a (SMP) control processor, any OpenMP compliant program can run on it anyway.

But the SMECY tools can use more or less information from available OpenMP decorations available in the code.

5.2 Mapping on hardware

The mapping of a function call on a specific piece of hardware can be specified with

```
#pragma smecy map(hardware[, unit])  
some_function_call(...);
```

- *hardware* is a symbol representing a hardware component of a given target such as CPU, GPP, GPU, PE... They are target specific.
- *unit* is an optional instance number for a specific hardware part. This is typically an integer starting a 0. This hardware number can be an expression of the environment to be able to have a loop managing different accelerators.

We can add an **if**(*scalar-expression*) to predicate hardware launching according some run-time expression to choose between hardware or local software execution, as in OpenMP with the same syntax. The idea is to be able to do a software execution if the data to process is too small compared to the latency of an hardware accelerator.

Recursion is not supported on hardware-mapped functions. If there are functions called from hardware-mapped functions, they will be automatically inlined (so no recursion allowed in them either). If a function is mapped to a more programmable hardware (GPP), recursions in these called functions should be allowed.

5.3 Producer/consumer information

To generate hardware communications where there is only a function call, the compiler need to figure out what is the memory zone used to execute the function and then what memory zone is written by the function *and* that will be used later in the program². From these information, copy-in and copy-out operations can be generated.

5.3.1 Function arguments

```
#pragma smecy arg(arg_id, arg_clause[, arg_clause]...)  
some_function_call(...);
```

Direction directive defines how the data flows between the caller and the function:

- **in** the argument is used by the function;
- **out** the argument is used by the function;
- **inout** the argument is used and produced by the function;
- **unused** the argument is not used and produced by the function.

Argument layout specifies how the argument is used in the function.

- An optional *array_size_descriptor* such as `[n] [m]` expressing that the data is used from the callee point of view as such an array starting at the address given in parameter. If not specified, all the caller argument is used;
- An optional *array_range_descriptor* restriction such as `/[n:m] [2] [3:7]` expressing that the data is used from the callee point of view as an array with only this element ranges used. If not specified, all the array is used according to its size specified or not. If only some ranges are lacking, all the matching dimension is used. For example `/[4] []` matches the column 4 of an array.

The more precise this description is and the less data transfers occur.

5.3.2 Global variables

Right now we do not deal with sharing information through global variables, because it is more difficult to track. Only function parameters are used to exchange information.

But we can imagine to map global variables with this clause:

```
#pragma smecy global_var(var, arg_clause[, arg_clause]...)
```

5.4 Remapping specification

TO FINISH

Since we always want a sequential equivalence, that means that the sequential code representing the computation on an accelerator really consume The easy

HPF

²If a function produces something not used later, it is useless to get it back.

5.5 Hardware specific pragma

To be defined in collaboration with the various hardware suppliers of the project (P2012, EdkDSP...).

6 SMECY high-level API

6.1 OpenMP

Since we support OpenMP pragma, we also support OpenMP API that allows for example:

- getting/setting the number of threads;
- getting the number of available processors on the current domain;
- manipulating locks.

6.2 MultiCore Association API

We rely on the MCAPI, MTAPI and MRAPI for low memory footprint light-weight communications, threading and synchronization. Refer to their documentations for more information.

There is a reference implementation based on Linux *pthread*s that can be used as an example to port to the various available hardware.

This can be used to express communicating process, asynchronous communications, synchronization, etc.

6.3 NPL API

NPL is the API defined to program ST P2012 in a native way. Since it is rather at the same level of MCAPI/MRAPI/MTAPI, it should be easy to implement one above the other. Since the MCA APIs are standard, we think that it is commercially interesting to provide a MCA APIs over NPL or other to widen P2012 usage.

Define here NPL

6.4 EdkDSP API

Define here what is useful in the project.

6.5 OpenCL

Since ST P2012 can be programmed in OpenCL which is also a programming API, a C process can use OpenCL orthogonally with other API. A kernel launching is done by defining the kernel source, the memory zone to use and to transfer and the different parameters of the kernel.

Refer to OpenCL documentation for more information.

7 Compilation

7.1 OpenMP support

From the programmer point of view it may be equivalent to have

- an OpenMP compiler generating SMECY API code such as the parallelism between SMECY target accelerators is run by OpenMP threads dealing each one with an hardware resource sequentially in a synchronous way;
- or a SMECY compiler understanding the OpenMP syntax and generating directly some parallel execution of SMECY accelerators in an asynchronous way.

7.2 Information available

7.2.1 Doxygen qualifier

In the Doxygen documentation mark-up language used in comments to detail various entities of a program, there are information such as **in**, **out** or **inout** qualifier on function parameters that can be used to generate the right communication with some hardware accelerators.

7.2.2 C qualifier

Qualifiers such as **const** attribute qualifier, address space names, register names, accumulator, saturation, etc. are used to generate the right target function or instruction.

7.2.3 Pragma

Of course the pragma information is heavily used in the compilation process

For example to generate correct communication, the mapping information is used, and from the used/defined information plus optional remapping information, a real communication with an API is used.

```
void
2 invert_vector(int line_size ,
                int input_line[line_size],
                int output_line[line_size]) {
    for(int i = 0; i < line_size; i++)
        output_line[i] = 500 - input_line[i];
7 }
    [...]
    int image[HEIGHT][WIDTH];
    [...]
12 #pragma smecy map(PE, proc & 7) \
    arg(2, in , [1][LINE_SIZE]) \
    arg(3, out , [1][LINE_SIZE])
    // Invert an horizontal line:
    invert_vector(LINE_SIZE,
17                 &image[HEIGHT - 20 - proc][WIDTH/2 + 2*proc],
                 &image[HEIGHT - 20 - proc][WIDTH/2 + 2*proc]);
```

can be compiled into

```

    int image[HEIGHT][WIDTH];
    /* First prepare the PE #(proc & 7) hardware to execute invert_vector.
3    That may be used to load some program or microcode, reconfigure a
        FPGA, load/compile an OpenCL kernel... */
    SMECY_set(PE, proc & 7, invert_vector);
    /* Send the WIDTH integer as arg 1 on invert_vector hardware function on
        PE #(proc & 7): */
8    SMECY_send_arg(PE, proc & 7, invert_vector, 1, int, LINE_SIZE);
    /* Send a vector of int of size LINE_SIZE as arg 2 on invert_vector
        hardware function on PE #(proc & 7): */
    SMECY_send_arg_vector(PE, proc & 7, invert_vector, 2, int,
        &image[HEIGHT - 20 - proc][WIDTH/2 + 2*proc], LINE_SIZE);
13 // Launch the hardware function or remote program:
    SMECY_launch(PE, 0, invert_vector);
    /* Get a vector of int of size LINE_SIZE as arg 3 on invert_vector
        hardware function on PE #(proc & 7): */
    SMECY_get_arg_vector(PE, proc & 7, invert_vector, 3, &image[HEIGHT - 20 - proc]
18                                [WIDTH/2 + 2*proc],
                                LINE_SIZE);

```

with low level macros described in § 8.

invert_vector() is either an already implemented function in a hardware library, or it is compiled by a target specific compiler with some callable interface.

For more complex calls needing remapping, such as:

```

    int input_line[LINE_SIZE];
    int output_line[LINE_SIZE];
    /* We need to remap data in the good shape. The compiler should use
        the remapping information to generate DMA transfer for example and
5    remove input_line array */
    SMECY_remap_int2D_to_int1D(HEIGHT, WIDTH, HEIGHT/3, 30 + 20*proc,
                                LINE_SIZE, 1, image,
                                LINE_SIZE, input_line);
    // Each iteration is on a different PE in parallel:
10 #pragma smecy map(PE, proc) arg(2, in, [LINE_SIZE]) arg(3, out, [LINE_SIZE])
    invert_vector(LINE_SIZE, input_line, output_line);
    SMECY_remap_int1D_to_int2D(LINE_SIZE, output_line,
                                HEIGHT, WIDTH, HEIGHT/3, 30 + 20*proc,
                                LINE_SIZE, 1, image);

```

is compiled by using other hardware interfaces involving more complex DMA:

```

1    // May not be useful if this function is already set:
    SMECY_set(PE, proc & 7, invert_vector);
    /* Send the WIDTH integer as arg 1 on invert_vector hardware function on
        PE #(proc & 7): */
    SMECY_send_arg(PE, proc & 7, invert_vector, 1, int, LINE_SIZE);
6    /* Send a vector of int of size LINE_SIZE as arg 2 on invert_vector
        hardware function on PE #(proc & 7) but read as a part of a 2D array: */
    smecy_send_arg_int_DMA_2D_PE_0_invert_vector(3, &image[HEIGHT - 20 - proc]
                                                    [WIDTH/2 + 2*proc],
                                                    HEIGHT, WIDTH, HEIGHT/3, 30 + 20*proc,
11                                LINE_SIZE, 1);
    SMECY_send_arg_DMA_2D_to_1D(PE, proc & 7, invert_vector, 2, int,
                                &image[HEIGHT - 20 - proc][WIDTH/2 + 2*proc],
                                HEIGHT, WIDTH, HEIGHT/3, 30 + 20*proc,

```

```

16                                     LINE_SIZE, 1);
// Launch the hardware function or remote program:
SMECY_launch(PE, 0, invert_vector);
SMECY_get_arg_DMA_1D_to_2D(PE, proc & 7, invert_vector, 3, int,
                           &image[HEIGHT - 20 - proc][WIDTH/2 + 2*proc],
                           HEIGHT, WIDTH, HEIGHT/3, 30 + 20*proc,
21                                     LINE_SIZE, 1);

```

7.3 Geometric inference

In the OpenMP SMP model, there is a global memory (well, with a weak coherence model) that may not exist in the execution model of a given target. For example, even if 2 hardware accelerators exchange information through memory according to the high-level programming model, in the real world we may have 2 processors communicating through message passing with MCAP on a NoC or 2 hardware accelerators connected through a pipeline.

To solve this issue, we use the OpenMP global memory like a scoreboard memory that is used to symbolically relate all the data flows and generates various communication schemes.

Since the memory dependencies are expressed by hyperparallelepipeds, we can do some intersection analysis to compute if a communication is needed or not between 2 devices of the target.

8 SMECY low level hardware API

To call real hardware accelerators, few C macros are needed to interface a program running on some processor to a function running on another processor or in some hardware accelerator.

Since the implementation may depend also on the processor calling the macros (not the same IO bus will be used from an x86 or a DSP to call the same operator), a global preprocessing symbol must be defined by the compiler before using these macros, such as

```
#define SMECY_LOCALPROC x86
```

or

```
#define SMECY_LOCALPROC DSP
```

Few macros are necessary:

```

/* Prepare a processing element to execute a function

   @param pe is the symbol of a processing element, such as GPP, DSP, PE...
4   @param[in] instance is the instance number of the processor element to use
   @param func is the function to load on the processor element
*/
#define SMECY_set(pe, instance, func) ...

9 /* Send a scalar argument to a function on a processing element

   @param pe is the symbol of a processing element, such as GPP, DSP, PE...
   @param[in] instance is the instance number of the processor element to use

```



```

14      @param func is the function to load on the processor element
      @param[in] arg is the argument instance to set
      @param type is the type of the scalar argument to send
      @param[in] val is the value of the argument to send
    */
19    #define SMECY_send_arg(pe, instance, func, arg, type, val) ...

    /* Send a vector argument to a function on a processing element

       @param pe is the symbol of a processing element, such as GPP, DSP, PE...
       @param[in] instance is the instance number of the processor element to use
24      @param func is the function to load on the processor element
       @param[in] arg is the argument instance to set
       @param type is the type of the vector element to send
       @param[in] addr is the starting address of the vector to read from
                           caller memory
29      @param[in] size is the length of the vector
    */
    #define SMECY_send_arg_vector(pe, instance, func, arg, type, addr, size) ...

    /* Launch the hardware function or remote program using previously loaded
34      arguments

       @param pe is the symbol of a processing element, such as GPP, DSP, PE...
       @param[in] instance is the instance number of the processor element to use
       @param func is the function to load on the processor element
39      A kernel can be launched several times without having to set/reset its function.
    */
    #define SMECY_launch(pe, instance, func) ...

44    /* Get the return value of a function on a processing element

       @param pe is the symbol of a processing element, such as GPP, DSP, PE...
       @param[in] instance is the instance number of the processor element to use
       @param func is the function to load on the processor element
49      @param type is the type of the scalar argument to send
       @return the value computed by the function
    */
    #define SMECY_get_return(pe, instance, func, type) ...

54    /* Get a vector value computed by a function on a processing element

       @param pe is the symbol of a processing element, such as GPP, DSP, PE...
       @param[in] instance is the instance number of the processor element to use
       @param func is the function to load on the processor element
59      @param[in] arg is the argument instance to retrieve
       @param type is the type of the vector element
       @param[out] addr is the starting address of the vector to write in
                           caller memory
       @param[in] size is the length of the vector
64    */
    #define SMECY_get_arg_vector(pe, instance, func, arg, type, addr, size) ...

```

```

/* Reset a processing element to execute a function

69    @param pe is the symbol of a processing element, such as GPP, DSP, PE...
    @param[in] instance is the instance number of the processor element to use
    @param func is the function to unload from the processor element

    This is used for example to remove consuming resources to decrease
74    power. Giving here the function name may be useful for weird case to
    avoid having short-circuit between CLB in a FPGA during unconfiguring stuff
*/
#define SMECY_reset(pe, instance, func) ...

```

There is also a macro for an asynchronous call and to wait for completion.